

CSE 573 - Project 2

Amlan Gupta (#50288686)

November 2018

1 Image Features and Homography

```
1 def perform_sift(img):
2
3     sift_img = np.zeros(img.shape)
4
5
6     sift = cv2.xfeatures2d.SIFT_create()
7
8     kp, desc = sift.detectAndCompute(img, None)
9     sift_img = cv2.drawKeypoints(img, kp, sift_img)
10
11     return sift_img, kp, desc
12
13 def find_knn_match(img_1, img_2, kp_1, kp_2, desc_1, desc_2):
14
15     bf = cv2.BFMatcher()
16     matches = bf.knnMatch(desc_1, desc_2, k=2)
17
18     # Apply ratio test
19     good_matches = []
20     pts1 = []
21     pts2 = []
22     for m, n in matches:
23         if m.distance < 0.75*n.distance:
24             good_matches.append(m)
25             pts2.append(kp_2[m.trainIdx].pt)
26             pts1.append(kp_1[m.queryIdx].pt)
27
28
29     knn_matched_img = np.zeros(img_1.shape)
30     knn_matched_img = cv2.drawMatches(img_1, kp_1, img_2, kp_2, good_matches, knn_matched_img, flags=2)
31     return knn_matched_img, good_matches, pts1, pts2
32
33
34 def find_homography_and_match_images(good_matches, kp_1, kp_2, img_1, img_2):
35
36     src_pts = np.float32([ kp_1[m.queryIdx].pt for m in good_matches ]).reshape((-1,1,2))
37     dst_pts = np.float32([ kp_2[m.trainIdx].pt for m in good_matches ]).reshape((-1,1,2))
38
39     M, mask = cv2.findHomography(src_pts, dst_pts, cv2.RANSAC, 5.0)
40     matchesMask = mask.ravel()
41     good_matches_np = np.asarray(good_matches)
42
43
44     matchesMask = matchesMask[matchesMask==1]
45     good_matches_np = good_matches_np[matchesMask==1]
46
47     randIdx = np.random.randint(low=0, high=good_matches_np.shape[0], size=10)
48
49     good_matches_np = good_matches_np[randIdx]
```

```

50     matchesMask = matchesMask[randIndx]
51
52
53     draw_params = dict(matchColor = (0, 0, 255),
54                         singlePointColor = None,
55                         matchesMask = matchesMask.tolist(),
56                         flags = 2)
57
58
59     matches_img = cv2.drawMatches(img_1, kp_1, img_2, kp_2, good_matches_np.tolist(), None, **draw_params)
60     return M, matches_img
61
62
63
64 def do_image_stitching(img_1, img_2, M):
65
66     (h1, w1) = img_1.shape[:2]
67     (h2, w2) = img_2.shape[:2]
68
69     #remap the coordinates of the projected image onto the panorama image space
70     top_left = np.dot(M,np.asarray([0,0,1]))
71     top_right = np.dot(M,np.asarray([w2,0,1]))
72     bottom_left = np.dot(M,np.asarray([0,h2,1]))
73     bottom_right = np.dot(M,np.asarray([w2,h2,1]))
74
75     #normalize
76     top_left = top_left/top_left[2]
77     top_right = top_right/top_right[2]
78     bottom_left = bottom_left/bottom_left[2]
79     bottom_right = bottom_right/bottom_right[2]
80
81
82
83     pano_left = int(min(top_left[0], bottom_left[0], 0))
84     pano_right = int(max(top_right[0], bottom_right[0], w1))
85     W = pano_right - pano_left
86
87     pano_top = int(min(top_left[1], top_right[1], 0))
88     pano_bottom = int(max(bottom_left[1], bottom_right[1], h1))
89     H = pano_bottom - pano_top
90
91     size = (W, H)
92
93     # offset of first image relative to panorama
94     X = int(min(top_left[0], bottom_left[0], 0))
95     Y = int(min(top_left[1], top_right[1], 0))
96     offset = (-X, -Y)
97
98     panorama = np.zeros((size[1], size[0]), np.uint8)
99
100     (ox, oy) = offset
101
102     translation = np.matrix([
103
104         [1.0, 0.0, ox],
105         [0, 1.0, oy],
106         [0.0, 0.0, 1.0]
107     ])
108
109     M = translation * M
110
111     cv2.warpPerspective(img_1, M, size, panorama)
112

```

```

113     panorama[oy:h1+oy, ox:ox+w1] = img_2
114
115     return panorama
116
117
118 def image_feature_and_homography(mountain_1_img, mountain_2_img):
119
120     # task 1.1
121
122     kp_mountain_1_img, kp_1, desc_1 = perform_sift(mountain_1_img)
123     write_image(kp_mountain_1_img, 'task1_sift1.jpg')
124
125
126     kp_mountain_2_img, kp_2, desc_2 = perform_sift(mountain_2_img)
127     write_image(kp_mountain_2_img, 'task1_sift2.jpg')
128
129
130     # task 1.2
131
132     knn_matched_img, good_matches, _, _ = find_knn_match(mountain_1_img, mountain_2_img,
133     kp_1, kp_2, desc_1, desc_2)
134     write_image(knn_matched_img, 'task1_matches_knn.jpg')
135
136
137     #task 1.3, task 1.4
138
139     h_matrix, task1_matches = find_homography_and_match_images(good_matches, kp_1, kp_2,
140     mountain_1_img, mountain_2_img)
141     print(h_matrix)
142     write_image(task1_matches, 'task1_matches.jpg')
143
144     # task 1.5
145     panorama = do_image_stitching(mountain_1_img, mountain_2_img, h_matrix)
146     write_image(panorama, 'task1_pano.jpg')
147
148
149
150
151 def main():
152
153     mountain_1_img = cv2.imread(SOURCE_FOLDER + "mountain1.jpg", 0)
154     mountain_2_img = cv2.imread(SOURCE_FOLDER + "mountain2.jpg", 0)
155
156     image_feature_and_homography(mountain_1_img, mountain_2_img)
157
158
159
160 main()

```

1. Given two images `mountain1.jpg` and `mountain2.jpg`, extract SIFT features and draw the keypoints for both images. Include the resulting two images (`task1_sift1.jpg`, `task1_sift2.jpg`) in the report.




Figure 1: Keypoints for `mountain1.jpg`




Figure 2: Keypoints for `mountain2.jpg`

- Match the keypoints using k-nearest neighbour ($k=2$), i.e., for a keypoint in the left image, finding the best 2 matches in the right image. Filter good matches satisfy $m.distance < 0.75 n.distance$, where m is the first match and n is the second match. Draw the match image using `cv2.drawMatches` for all matches (your match image should contain both inliers and outliers). Include the result image (task1_matches_knn.jpg) in the report.




Figure 3: Matched keypoints using k-nearest neighbour ($k=2$)

- Compute the homography matrix H (with RANSAC) from the first image to the second image. Include the matrix values in the report.

$$\begin{bmatrix} 1.589302 & -0.291559 & -395.969265 \\ 0.449424 & 1.431109 & -190.613988 \\ 0.001213 & -0.000063 & 1.000000 \end{bmatrix}$$

- Draw the match image for around 10 random matches using only inliers. Include the result image (task1_matches.jpg) in the report.




Figure 4: 10 random matches using only inliers

5. Warp the first image to the second image using H . The resulting image should contain all pixels in mountain1.jpg and mountain2.jpg. Include the result image (task1_pano.jpg) in the report.




Figure 5: Warped image using Homography matrix

2 Epipolar Geometry

```
1 def perform_sift(img):
2
3     sift_img = np.zeros(img.shape)
4
5
6     sift = cv2.xfeatures2d.SIFT_create()
7
8     kp, desc = sift.detectAndCompute(img, None)
9     sift_img = cv2.drawKeypoints(img, kp, sift_img)
10
11     return sift_img, kp, desc
12
13 def find_knn_match(img_1, img_2, kp_1, kp_2, desc_1, desc_2):
14
15     bf = cv2.BFMatcher()
16     matches = bf.knnMatch(desc_1, desc_2, k=2)
17
18     # Apply ratio test
19     good_matches = []
20     pts1 = []
21     pts2 = []
22     for m, n in matches:
23         if m.distance < 0.75*n.distance:
24             good_matches.append(m)
25             pts2.append(kp_2[m.trainIdx].pt)
26             pts1.append(kp_1[m.queryIdx].pt)
27
28
29     knn_matched_img = np.zeros(img_1.shape)
30     knn_matched_img = cv2.drawMatches(img_1, kp_1, img_2, kp_2, good_matches, knn_matched_img, flags=2)
31     return knn_matched_img, good_matches, pts1, pts2
32
```

```

33
34
35 def drawlines(img1,img2,lines,pts1,pts2, color):
36     r,c = img1.shape
37     img1 = cv2.cvtColor(img1,cv2.COLOR_GRAY2BGR)
38     img2 = cv2.cvtColor(img2,cv2.COLOR_GRAY2BGR)
39     for r,pt1,pt2,colr in zip(lines,pts1,pts2,color):
40         x0,y0 = map(int, [0, -r[2]/r[1] ])
41         x1,y1 = map(int, [c, -(r[2]+r[0]*c)/r[1] ])
42         img1 = cv2.line(img1, (x0,y0), (x1,y1), tuple(colr),1)
43         img1 = cv2.circle(img1,tuple(pt1),5,tuple(colr),-1)
44         img2 = cv2.circle(img2,tuple(pt2),5,tuple(colr),-1)
45     return img1,img2
46
47 def find_disparity_map(imgL, imgR):
48
49     window_size = 3
50     min_disp = 16
51     num_disp = 64-min_disp
52     stereo = cv2.StereoSGBM_create(minDisparity = min_disp,
53         numDisparities = num_disp,
54         blockSize = 9,
55         P1 = 8*3*window_size**2,
56         P2 = 32*3*window_size**2,
57         disp12MaxDiff = 1,
58         uniquenessRatio = 10,
59         speckleWindowSize = 100,
60         speckleRange = 32
61     )
62
63     disp = stereo.compute(imgL, imgR).astype(np.float32) / 16.0
64     disp = (disp-min_disp)/num_disp
65
66     disp = disp*300
67
68
69     write_image(disp,'task2_disparity.jpg')
70
71
72 def epipolar_geometry(tsucuba_left_img, tsucuba_right_img):
73
74     # task 2.1
75
76     kp_tsucuba_left_img, kp_1, desc_1 = perform_sift(tsucuba_left_img)
77     write_image(kp_tsucuba_left_img, 'task2_sift1.jpg')
78
79     kp_tsucuba_right_img, kp_2, desc_2 = perform_sift(tsucuba_right_img)
80     write_image(kp_tsucuba_right_img, 'task2_sift2.jpg')
81
82     # task 2.2, 2.3
83
84     knn_matched_img, good_matches, pts1, pts2 = find_knn_match(tsucuba_left_img, tsucuba_right_img,
85         kp_1, kp_2, desc_1, desc_2)
86
87     pts1 = np.int32(pts1)
88     pts2 = np.int32(pts2)
89     F, mask = cv2.findFundamentalMat(pts1,pts2,cv2.FM_RANSAC)
90     pts1 = pts1[mask.ravel()==1]
91     pts2 = pts2[mask.ravel()==1]
92     print(F)
93
94
95     randIndx = np.random.randint(low=0, high=pts1.shape[0], size=10)

```

```


96     pts1 = pts1[randIndx]
97     pts2 = pts2[randIndx]
98
99
100     color = np.random.randint(0,255, size=(10, 3)).tolist()
101
102     lines1 = cv2.computeCorrespondEpilines(pts2.reshape(-1,1,2), 2, F)
103     lines1 = lines1.reshape(-1,3)
104     tsucuba_left_ep, _ = drawlines(tsucuba_left_img, tsucuba_right_img, lines1, pts1, pts2, color)
105
106     write_image(tsucuba_left_ep, 'task2_epi_left.jpg')
107
108     lines2 = cv2.computeCorrespondEpilines(pts1.reshape(-1,1,2), 1, F)
109     lines2 = lines2.reshape(-1,3)
110     tsucuba_right_ep, _ = drawlines(tsucuba_right_img, tsucuba_left_img, lines2, pts2, pts1, color)
111
112     write_image(tsucuba_right_ep, 'task2_epi_right.jpg')
113
114     # task 2.4
115
116     find_disparity_map(tsucuba_left_img, tsucuba_right_img)
117
118
119
120
121 def main():
122
123     tsucuba_left_img = cv2.imread(SOURCE_FOLDER + "tsucuba_left.png", 0)
124     tsucuba_right_img = cv2.imread(SOURCE_FOLDER + "tsucuba_right.png", 0)
125
126     epipolar_geometry(tsucuba_left_img, tsucuba_right_img)
127
128
129
130
131 main()

```

1. Given two images `tsucuba_left.png` and `tsucuba_right.png`, do the same process for Task 1.1 and 1.2. Include the three images (2 for task 1.1 and 1 for task 1.2) (`task2_sift1.jpg`, `task2_sift2.jpg`, `task2_matches_knn.jpg`) in the report.



(a) for `tsucuba_left.png`



(b) for `tsucuba_right.png`

Figure 6: Detected Keypoints




Figure 7: Matched keypoints using k-nearest neighbour ($k=2$)


2. Computer the fundamental matrix F (with RANSAC). Include the matrix values in the report.

$$\begin{bmatrix} 3.88435667e^{-06} & -2.28e^{-04} & 4.41240231e^{-02} \\ 2.42e^{-04} & -1.90e^{-05} & 9.95e^{-02} \\ -4.58e^{-02} & -1.02e^{-01} & 1 \end{bmatrix}$$

3. Randomly select 10 inlier match pairs. For each keypoint in the left image, compute the epilines and draw on the right image. For each keypoint in the right image, compute the epilines and draw on the left image [Using different colors for different match pairs, but the same color for epilines on the left and right images with the same match pair.] Include two images (task2_epi_right.jpg, task2_epi_left.jpg) with epilines in the report.



(a) task2_epi_right.jpg



(b) task2_epi_left.jpg

Figure 8: Epilines

4. Compute the disparity map for tsucuba_left.png and tsucuba_right.png. Include the disparity image (task2_disparity.jpg) in the report.




Figure 9: Disparity map for tsucuba_left.png and tsucuba_right.png

3 K-means Clustering

```
1 def measure_euclidean_distance(pt1 , pt2):
2     dis = math.sqrt(((pt1[0] - pt2[0])**2) + ((pt1[1] - pt2[1])**2))
3     return dis
4
5 def calculate_distances_from_centroids(mu, mu_c, X, it_n):
6     cluster_c = []
7
8     for pt in X:
9
10         isFirst = True
11         for m, mc in zip(mu, mu_c):
12             if isFirst == True:
13                 d = measure_euclidean_distance(pt, m)
14                 c = mc
15                 isFirst = False
16             elif (measure_euclidean_distance(pt, m) < d):
17                 d = measure_euclidean_distance(pt, m)
18                 c = mc
19
20         cluster_c.append(c)
21
22     fig = plt.figure()
23     ax = fig.add_subplot(111)
24     plt.scatter(X[:,0], X[:,1], marker="^", facecolors="None", edgecolors= cluster_c)
25     plt.scatter(mu[:,0], mu[:,1], c= mu_c)
26     for xy in zip(mu[:,0], mu[:,1]):
27         ax.annotate('(%0.2f, %0.2f)' % xy, xy=xy, textcoords='data')
28     # plt.show()
29     plt.savefig(OUTPUT_FOLDER + 'task3_iter'+str(it_n+1)+'_a.jpg')
30     plt.clf()
31
32     return np.asarray(cluster_c)
33
34 def vector_classification():
35
36     X = np.array([
37         [5.9, 3.2],
38         [4.6, 2.9],
```

```

39     [6.2, 2.8],
40     [4.7, 3.2],
41     [5.5, 4.2],
42     [5.0, 3.0],
43     [4.9, 3.1],
44     [6.7, 3.1],
45     [5.1, 3.8],
46     [6.0, 3.0]])
47
48 mu = np.array([[6.2, 3.2], [6.6, 3.7], [6.5, 3.0]])
49 mu_c = ['r', 'g', 'b']
50
51 for i in range(2):
52
53     cluster_c = calculate_distances_from_centroids(mu, mu_c, X, i)
54
55     print(cluster_c)
56
57     clusters = []
58     for mc in zip(mu_c):
59         clusters.append(X[cluster_c == mc])
60
61     mu = []
62     for clus in clusters:
63         mu.append(np.mean(clus, axis=0))
64
65     mu = np.asarray(mu)
66     print(mu)
67
68
69     fig = plt.figure()
70     ax = fig.add_subplot(111)
71     plt.scatter(X[:,0], X[:,1], marker= "^", facecolors="None", edgecolors= cluster_c)
72     plt.scatter(mu[:,0], mu[:,1], c= mu_c)
73     for xy in zip(mu[:,0], mu[:,1]):
74         ax.annotate('({:.2f}, {:.2f})'.format(xy[0], xy[1]), xy=xy, textcoords='data')
75     plt.savefig(OUTPUT_FOLDER + 'task3_iter'+str(i+1)+'_b.jpg')
76     plt.clf()
77
78
79
80 def measure_euclidean_distance_3d(color1 , color2):
81     try:
82         dis = ((color1[0] - color2[0])**2) + ((color1[1] - color2[1])**2) + ((color1[2] - color2[2])**2)
83     except:
84         print(color1, color2)
85     return dis
86
87
88 def calculate_distances_from_centroids_3d(mu, mu_c, image):
89     cluster_c = np.zeros([image.shape[0], image.shape[1]])
90
91     h, w, l = image.shape
92
93
94     for i in range(h):
95         for j in range(w):
96             pixel = image[i][j]
97
98             isFirst = True
99             for m, mc in zip(mu, mu_c):
100                 if isFirst == True:
101                     d = measure_euclidean_distance_3d(pixel, m)

```

```

102         c = mc
103         isFirst = False
104         elif (measure_euclidean_distance_3d(pixel,m) < d):
105             d = measure_euclidean_distance_3d(pixel,m)
106             c = mc
107
108         cluster_c[i][j] = c
109
110
111     return np.asarray(cluster_c)
112
113
114 def color_quantization(image):
115
116
117     k = [3, 5, 10, 20]
118     # k = [20]
119
120     for kval in k:
121
122         mu = np.random.randint(0,255, size=(kval, 3))
123         mu_indx = np.random.randint(0,image.shape[0], size=(kval, 2))
124         mu = []
125
126         for i in range(mu_indx.shape[0]):
127             mu.append(image[mu_indx[i][0]][mu_indx[i][1]])
128         mu = np.asarray(mu).astype(float)
129
130         mu_c = np.arange(kval)
131
132
133         for zz in range(30):
134
135             cluster_c = calculate_distances_from_centroids_3d(mu, mu_c, image)
136
137             h, w, l = image.shape
138             clusters = []
139
140
141             for mc in zip(mu_c):
142                 clustered_img_np = []
143                 for i in range(h):
144                     for j in range(w):
145                         if(cluster_c[i][j] == mc):
146                             clustered_img_np.append(image[i][j])
147                 clusters.append(np.asarray(clustered_img_np))
148
149             mu = []
150             for clus in clusters:
151                 c_mean = np.nanmean(clus, axis=0)
152                 mu.append(c_mean)
153
154             mu = np.asarray(mu)
155
156             h, w, l = image.shape
157             output_img = np.zeros([h,w,l])
158
159             for i in range(h):
160                 for j in range(w):
161                     index =int(cluster_c[i][j])
162                     output_img[i][j] = mu[index]
163
164             op = output_img.astype(int)

```

```

165         print(zz)
166         write_image(op, 'task3_baboon_'+ str(kval) + '.jpg')
167         # print('baboon for ' + str(kval) + 'k means generated!')
168
169
170
171 def k_means_clustering(image):
172
173     #task 3.1, 3.2, 3.3
174     vector_classification()
175
176     #task 3.4
177     color_quantization(image)
178
179
180
181 def main():
182
183
184     baboon_img = cv2.imread(SOURCE_FOLDER + "baboon.jpg")
185     k_means_clustering(baboon_img)
186
187     print('DONE!!')
188
189
190
191
192 main()

```

1. Classify $N = 10$ samples according to nearest μ_i ($i = 1, 2, 3$). Plot the results by coloring the empty triangles in red, blue or green. Include the classification vector and the classification plot (task3_iter1_a.jpg) in the report.

The classification vector is [r r b r g r r b r r]




Figure 10: Classification plot for $N = 10$ samples

2. Recompute μ_i . Plot the updated μ_i in solid circle in red, blue, and green respectively. Include the updated μ_i values and the plot in the report (task3_iter1_b.jpg).

The value of updated μ_i in first iteration is:

$$\begin{bmatrix} 5.17142857 & 3.17142857 \\ 5.5 & 4.2 \\ 6.45 & 2.95 \end{bmatrix}$$




Figure 11: Updated μ_i in plot

3. For a second iteration, plot the classification plot and updated μ_i plot for the second iteration. Include the classification vector and updated μ_i values and these two plots (task3_iter2_a.jpg, task3_iter2_b.jpg) in the report.

Classification vector for second iteration: [b r b r g r r b g b]

The value of updated μ_i in second iteration is:

$$\begin{bmatrix} 4.8 & 3.05 \\ 5.3 & 4.0 \\ 6.2 & 3.025 \end{bmatrix}$$




Figure 12: task3_iter2_a.jpg




Figure 13: task3_iter2.b.jpg

4. Apply k-means to image color quantization. Using only k colors to represent the image baboon.jpg. Include the color quantized images for $k = 3, 5, 10, 20$ (task3_baboon_3.jpg, task3_baboon_5.jpg, task3_baboon_10.jpg, task3_baboon_20.jpg)



Figure 14: Color quantized baboon image for $k = 3$



Figure 15: Color quantized baboon image for $k=5$



Figure 16: Color quantized baboon image for $k=10$




Figure 17: Color quantied baboon image for $k = 20$